
Introduction to Recommendation Engine

— Jeff Liu —
2018/05/12

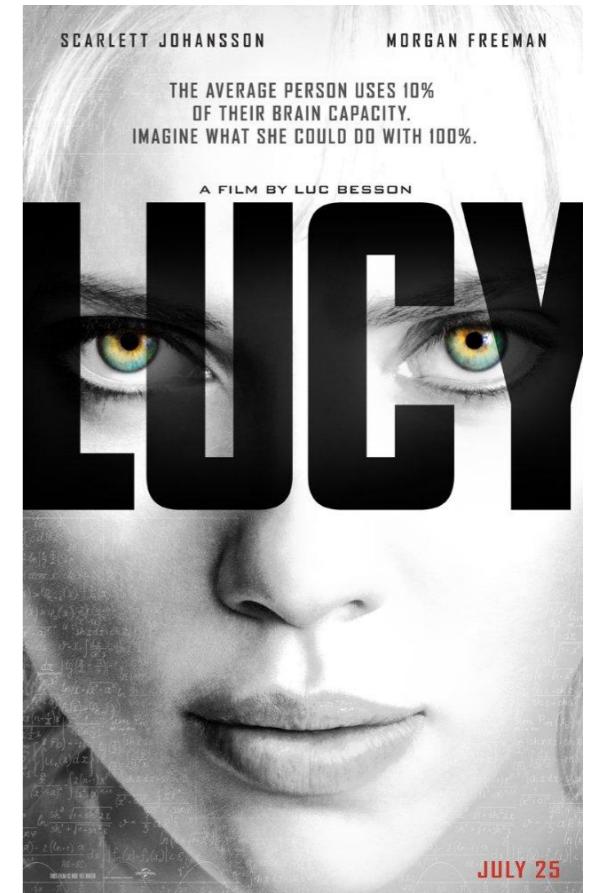
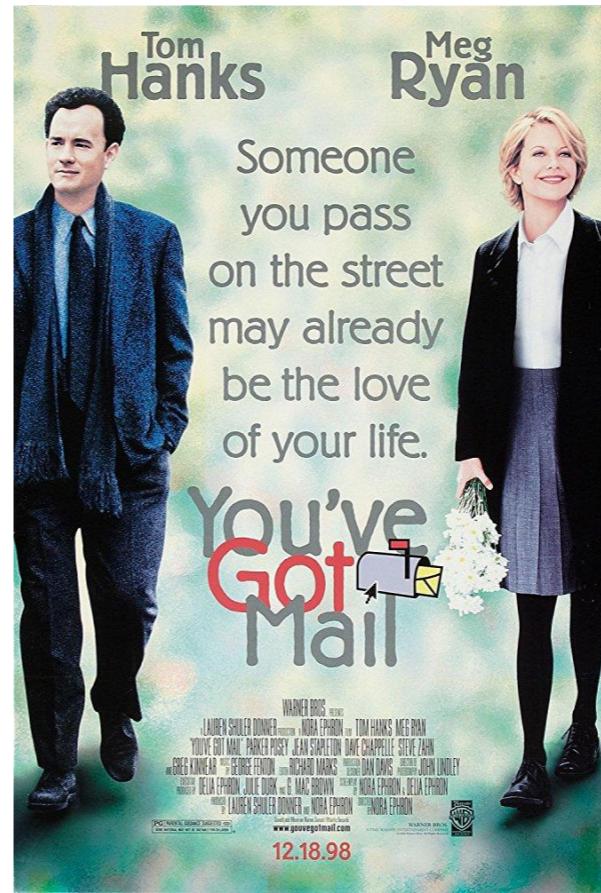
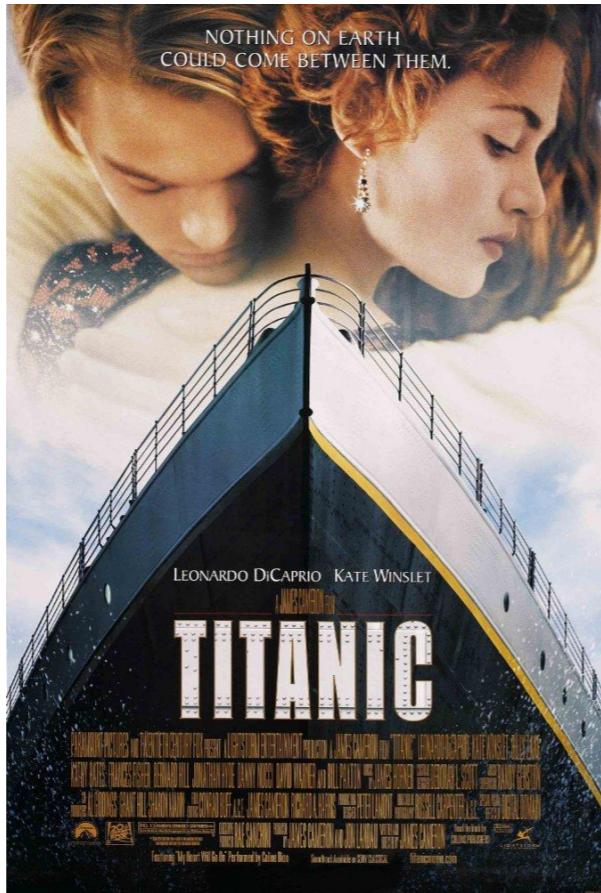
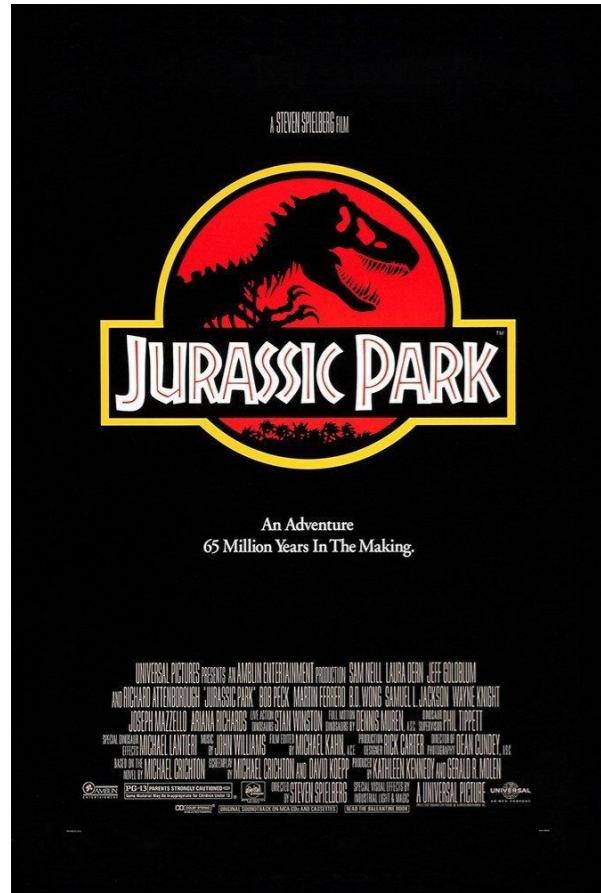
Prologue

- People are increasing the reliance on conveniences such as e-commerce store or streaming entertainment.
- “Guess” what the customers may like in advance, so to promote more things to sell and in turn generate more revenue.
- A recommendation engine (RE) is any kind of model that can infer the relationship between users/items and make proper prediction for the users.





Gender	Male
Age	30
Prefer genre	Sci-fi, comedy, action
Prefer director	Steven Spielberg, Christopher Nolan, Michael Bay, James Cameron
Prefer actor/actress	Tom Hanks, Leonardo DiCaprio, Anne Hathaway, Scarlet Johansson

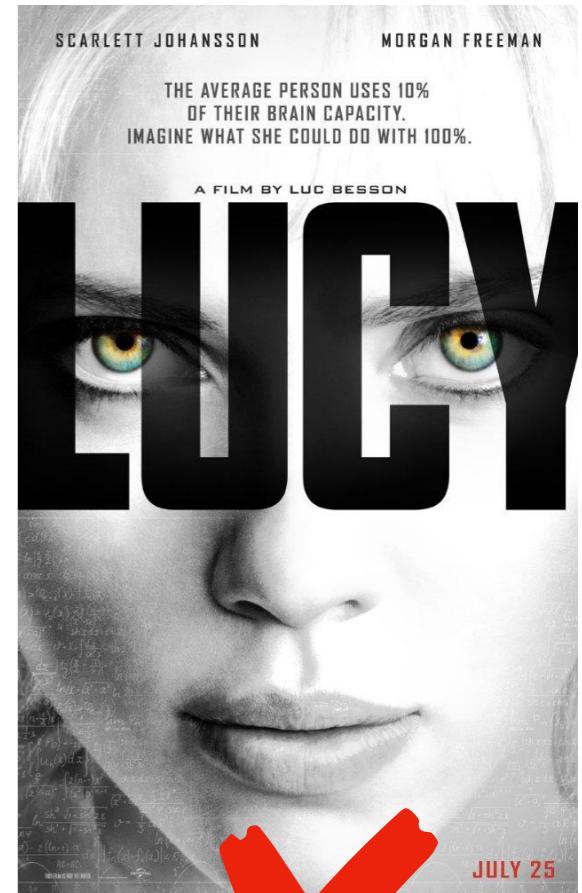
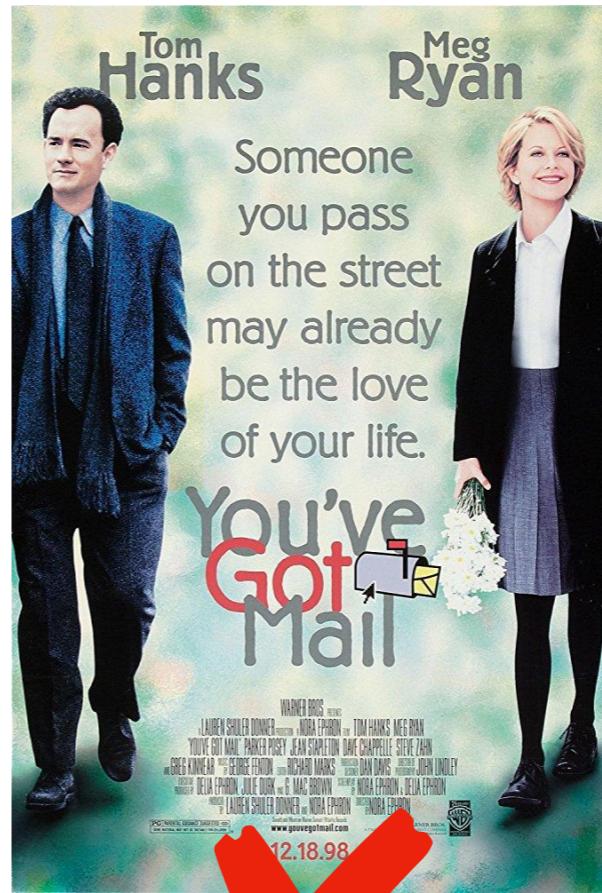
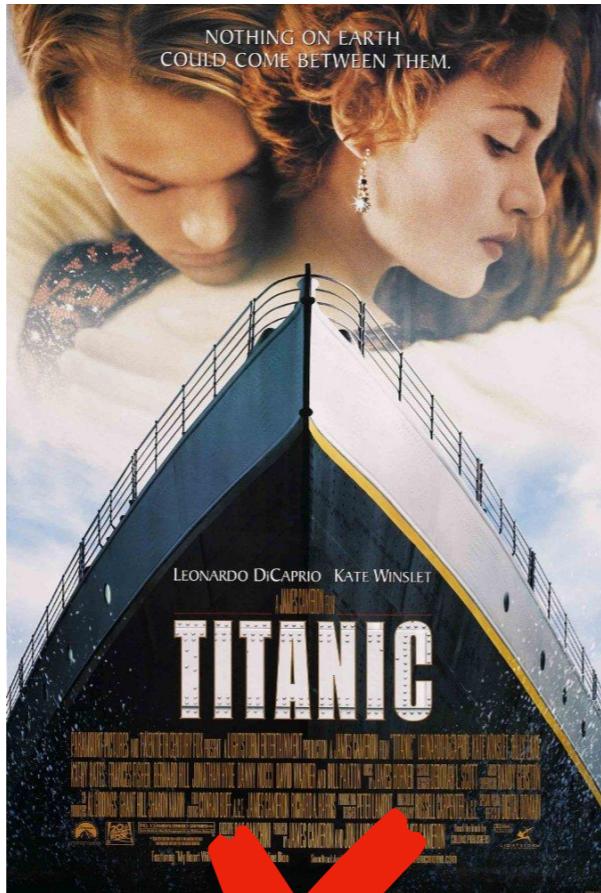
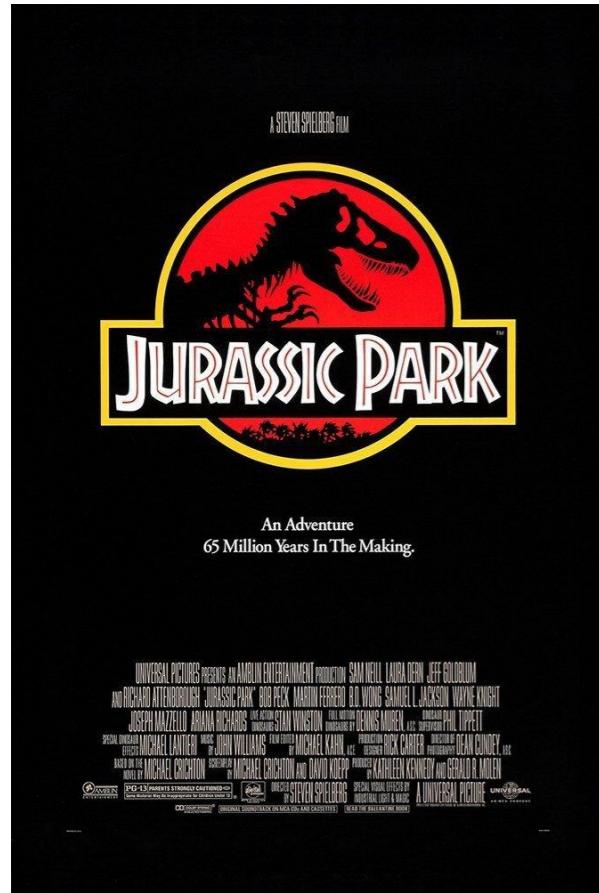


Content Filtering

- What we have just conducted is essentially one way of doing recommendation: *content filtering*.
- By building profiles for both the users and movies, we can provide recommendation by matching the *content* between the two groups, hence the name.
- Take a lot of efforts to build such profiles and many conditional settings to fit a person's taste. 



Gender	Male
Age	30
Prefer genre	Sci-fi, comedy, action
Prefer director	Steven Spielberg, Christopher Nolan, Michael Bay, James Cameron
Prefer actor/actress	Tom Hanks, Leonardo DiCaprio, Anne Hathaway, Scarlet Johansson



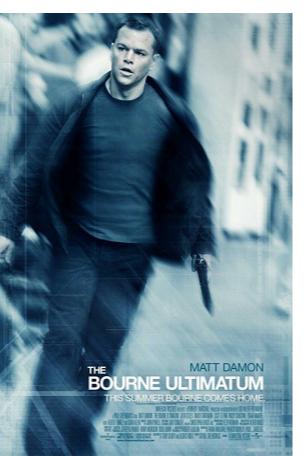
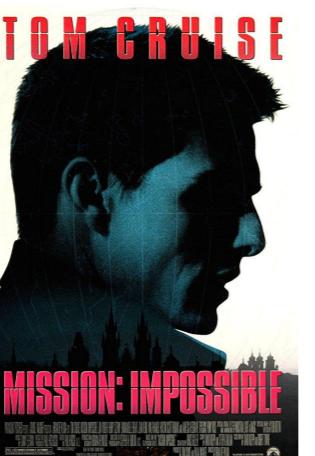
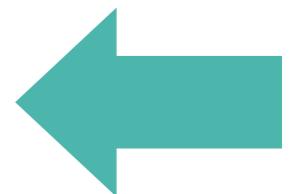
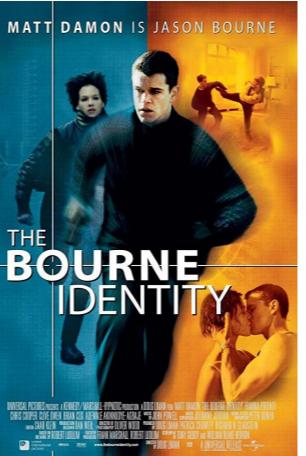
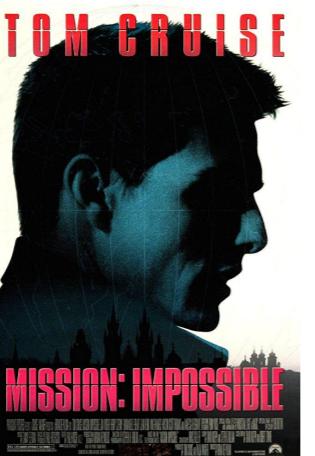
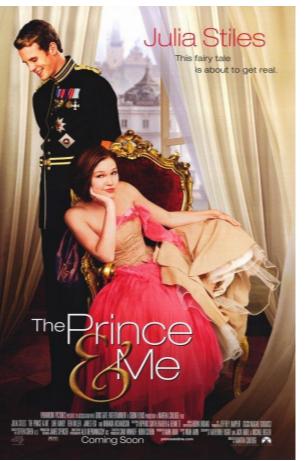
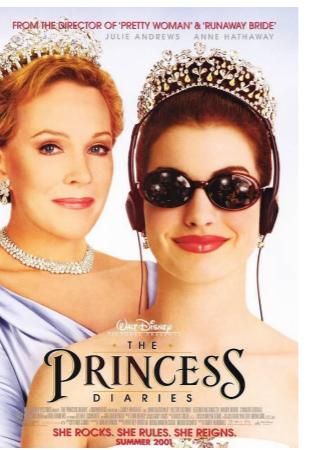
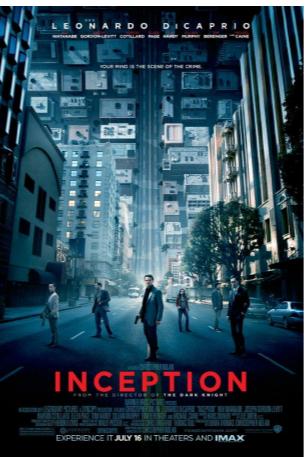
X

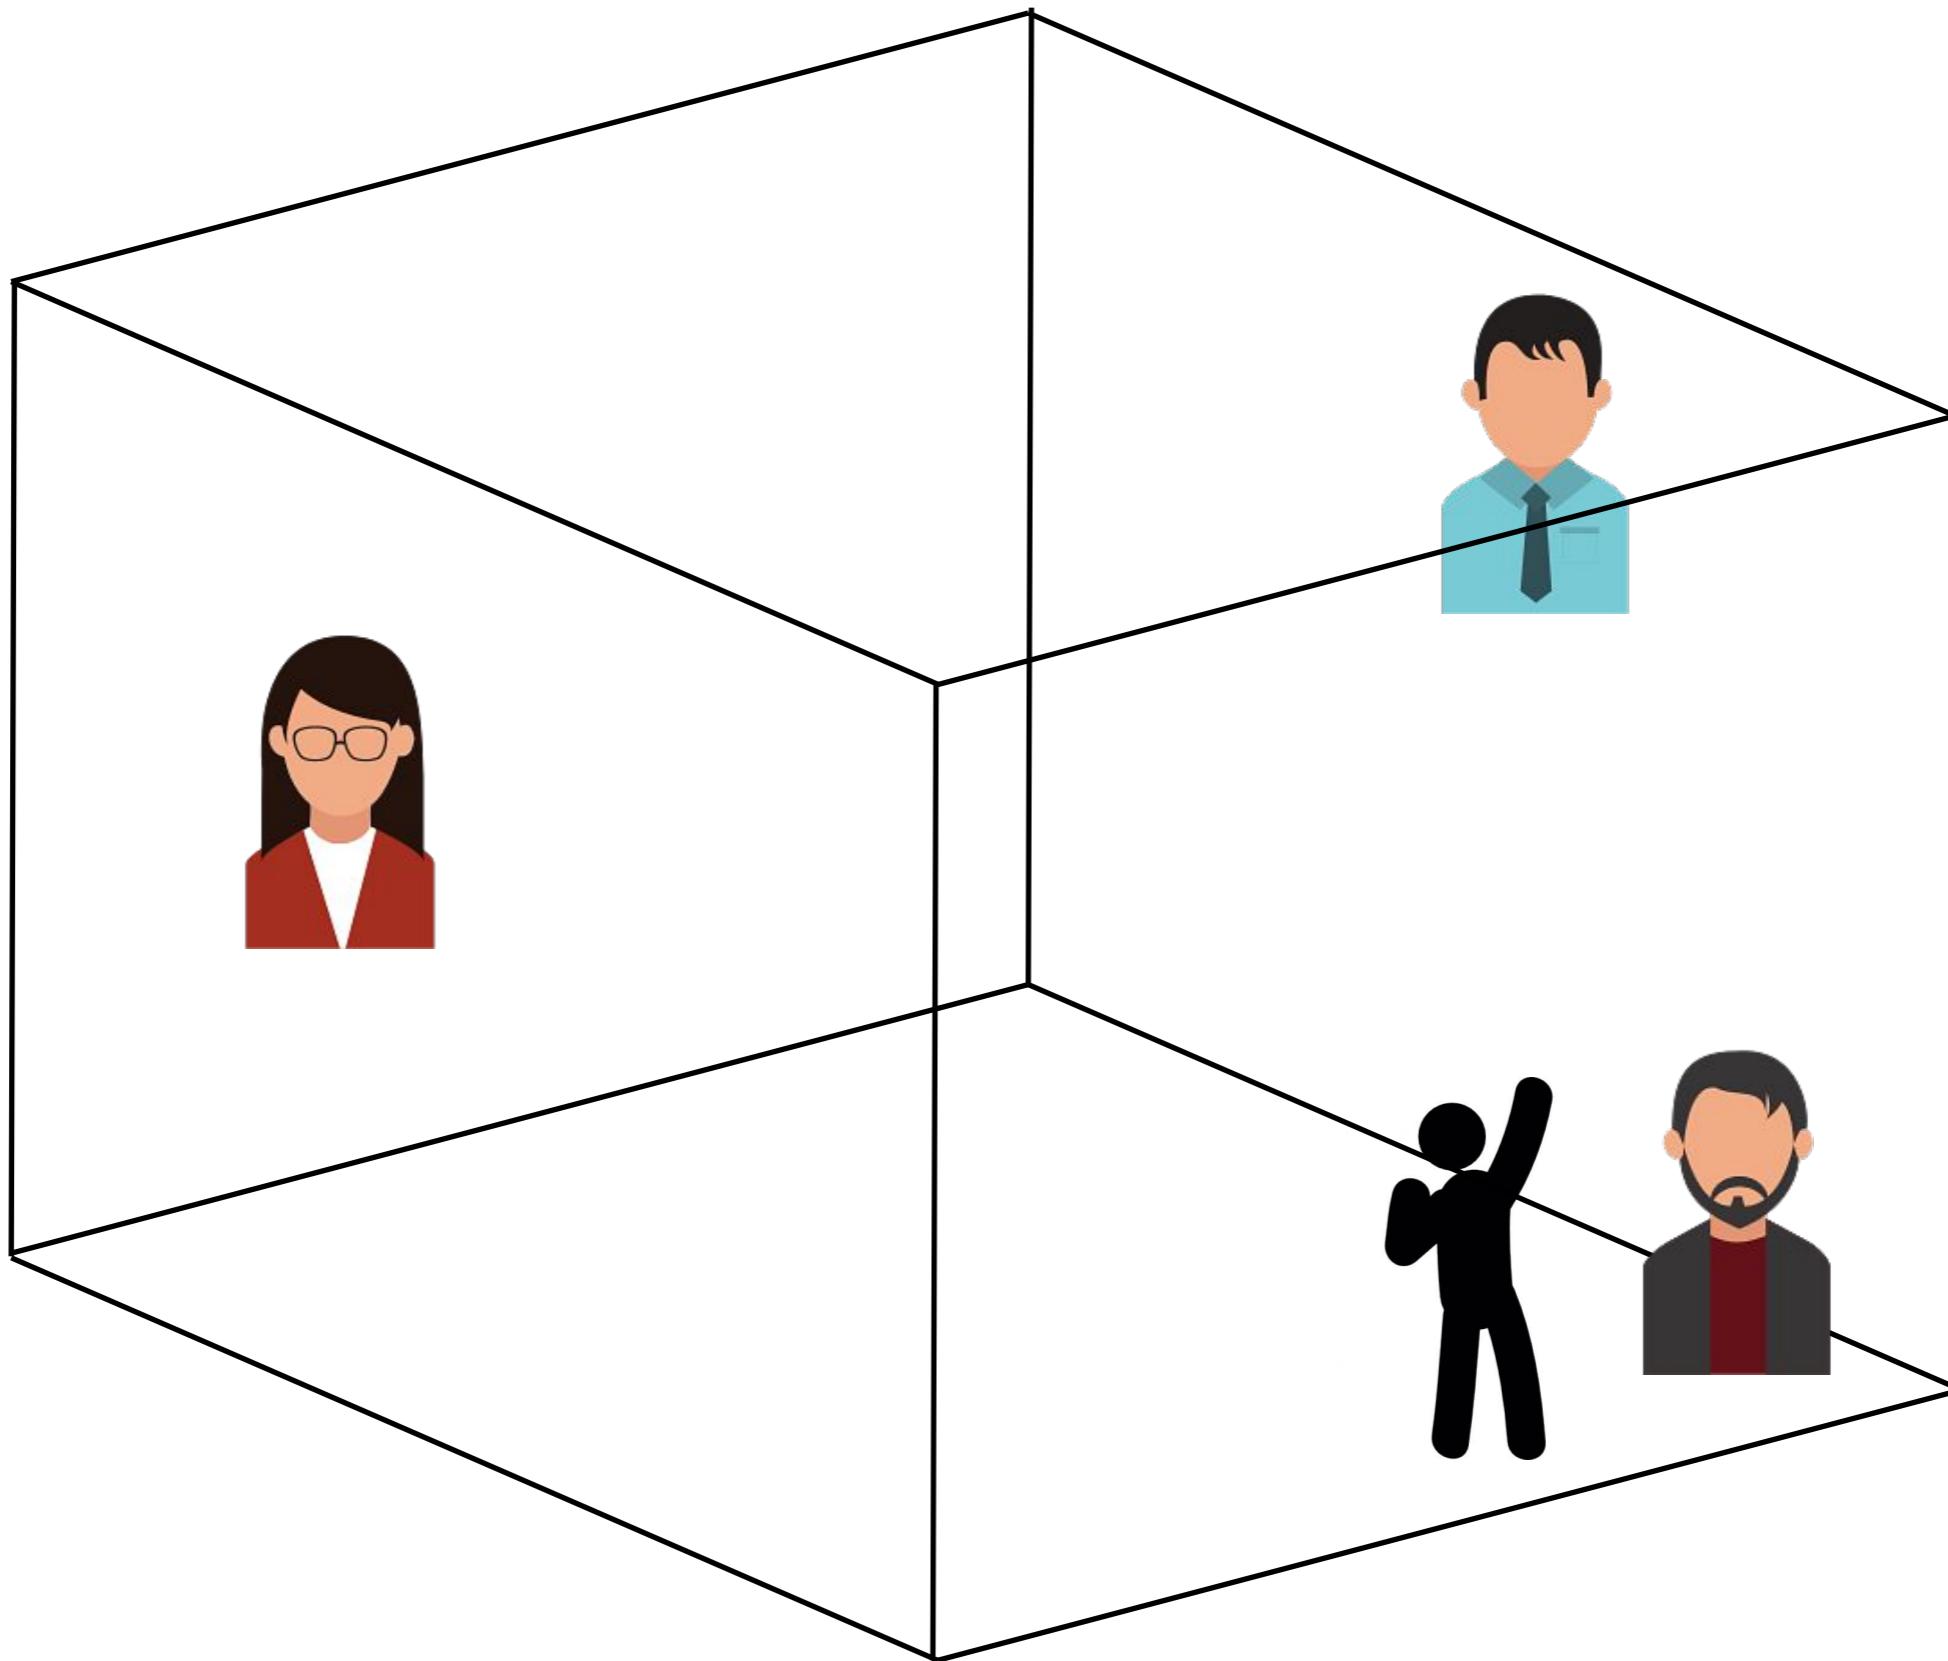
X

X

Strategies for Recommendation

- Beside *content filtering* we had just mentioned, there is another method called *collaborative filtering (CF)*.
- *CF* relies on past user behavior among a group of users (hence *collaborative*), so we aren't required to create profiles explicitly.
- The two primary areas of *CF* are the *neighborhood methods (memory-based)* and *latent factor models*.





User-based and Item-based CF

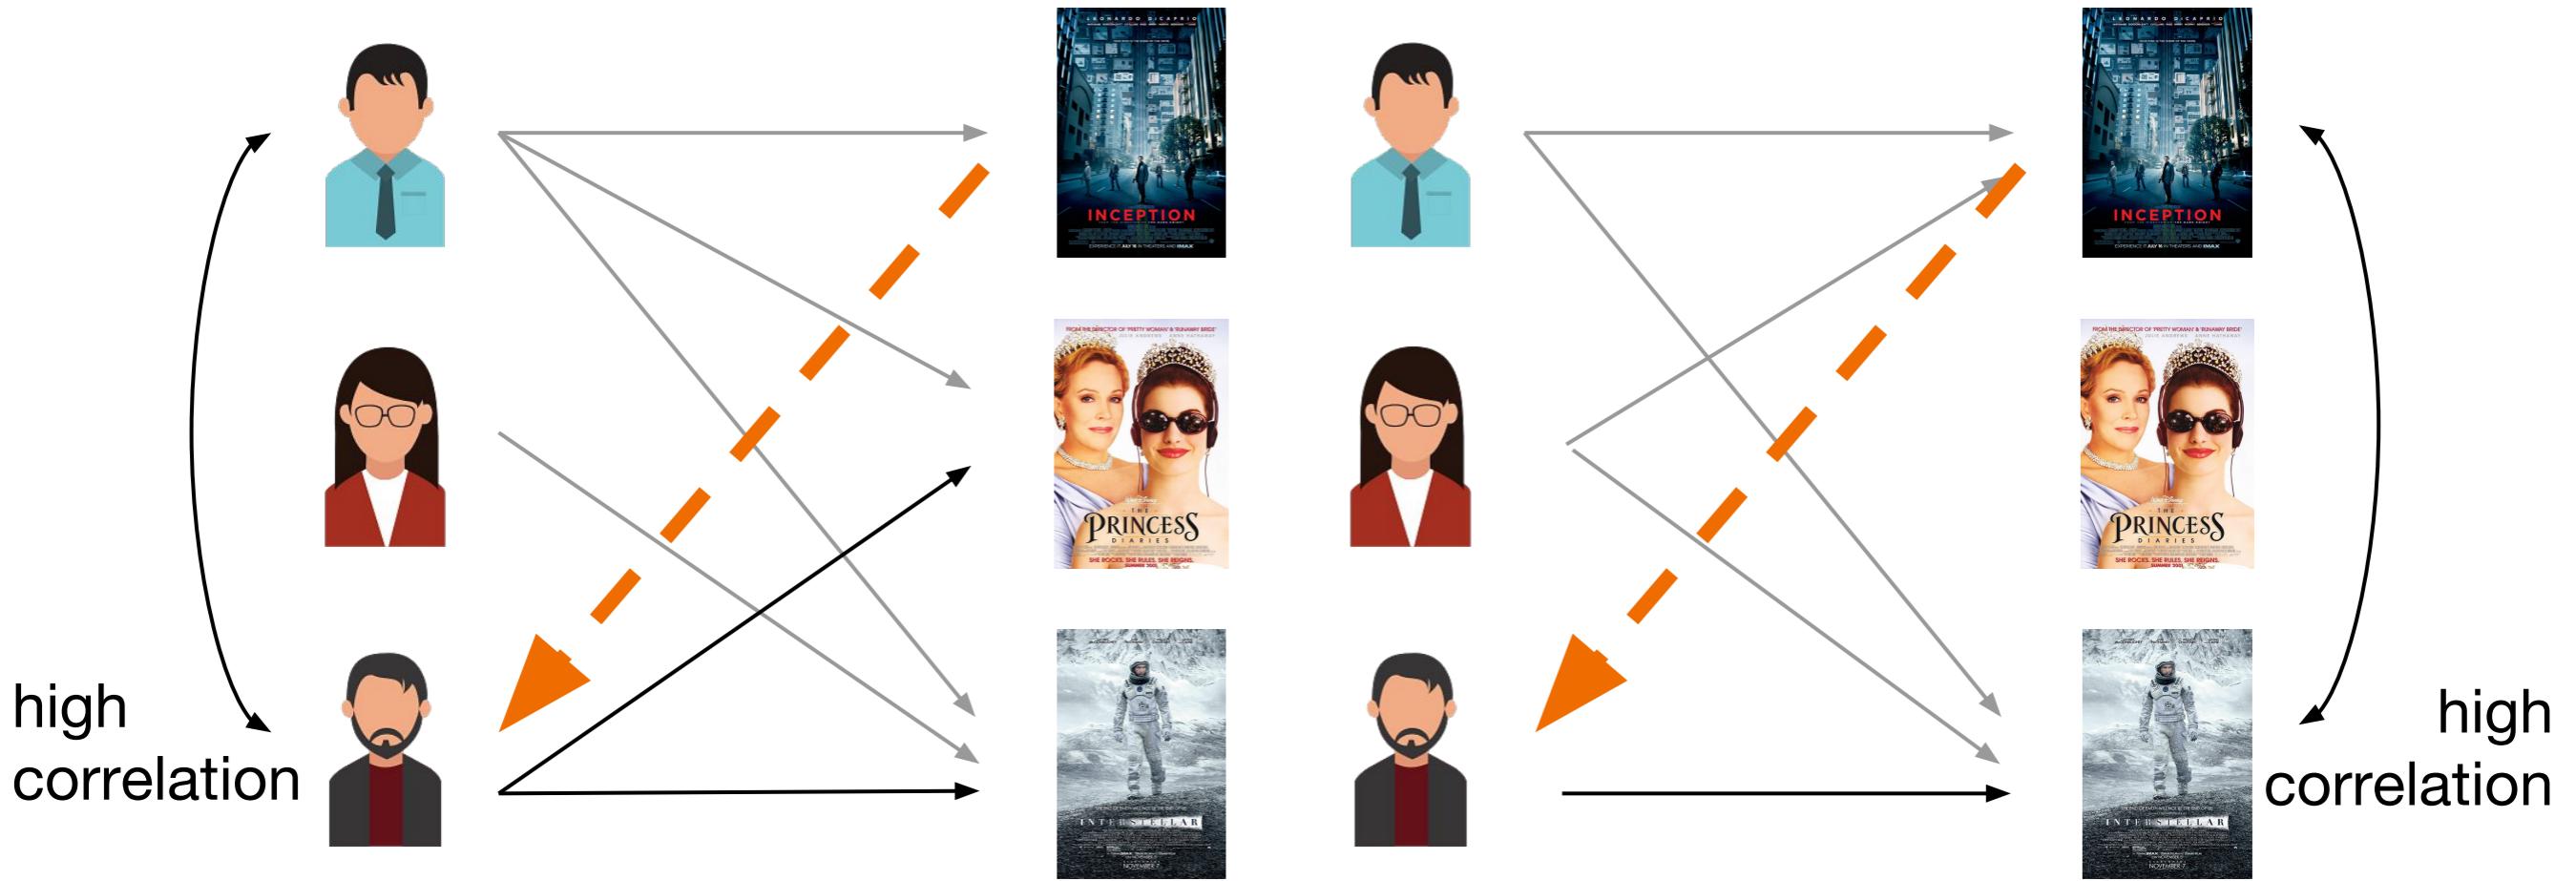
- For neighborhood methods, we have two primary types of algorithms: **user-based** and **item-based**.
- In *user-based CF*, we group together users who gave similar ratings to the same set of items, whereby we could later use the ratings of a specific user to predict those of his/her peers.
- For *item-based CF*, we group the items instead.



8.7	10	10	9	?	?	?	9	8	8.5
?	?	?	1	10	10	10	?	?	2
●									
●									
●									
●									
8	9	9	9	?	?	?	10	10	9.5

--- User-based CF

--- Item-based CF



User-based filtering

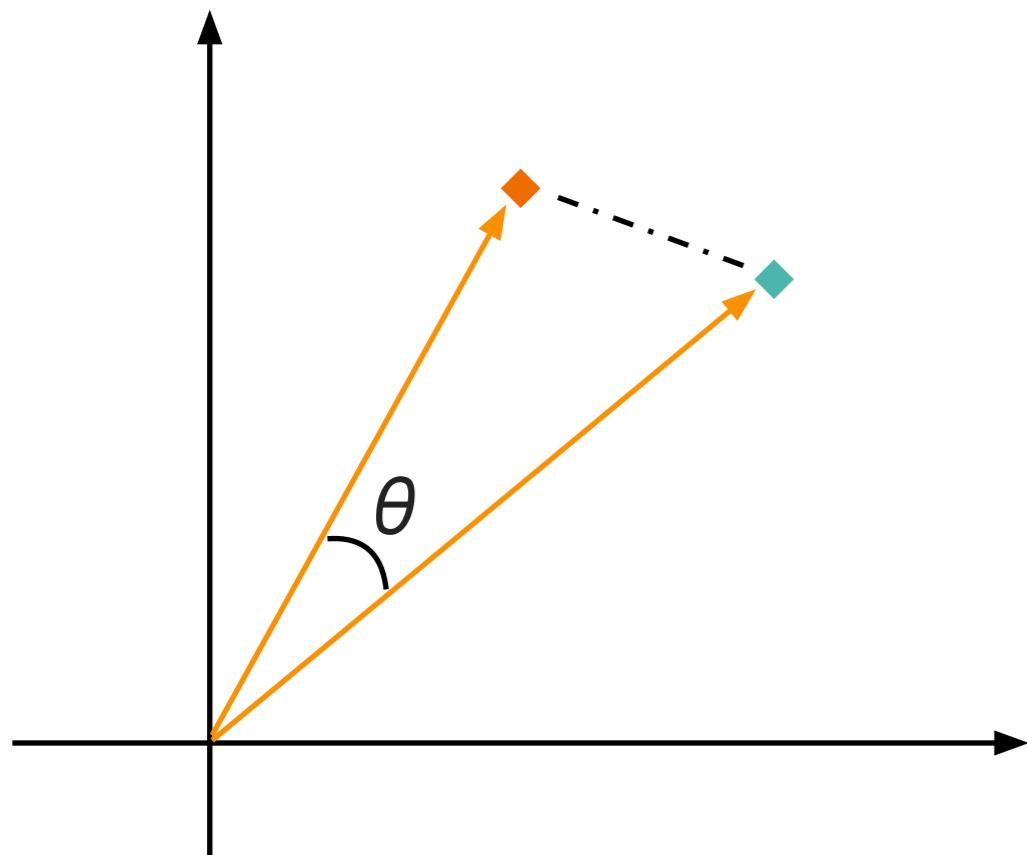
You may like it because your “friends” liked it.

Item-based filtering

You tend to like that item since you have liked those items.

Measuring Similarity

- To measure the similarity between users or items, we can use metrics like **Euclidean distance** or **cosine similarity**.



Euclidean distance: distance between points

$$d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2}$$

Cosine similarity: angle between vectors

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|_2 \|\mathbf{B}\|_2} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Example: User-Based CF



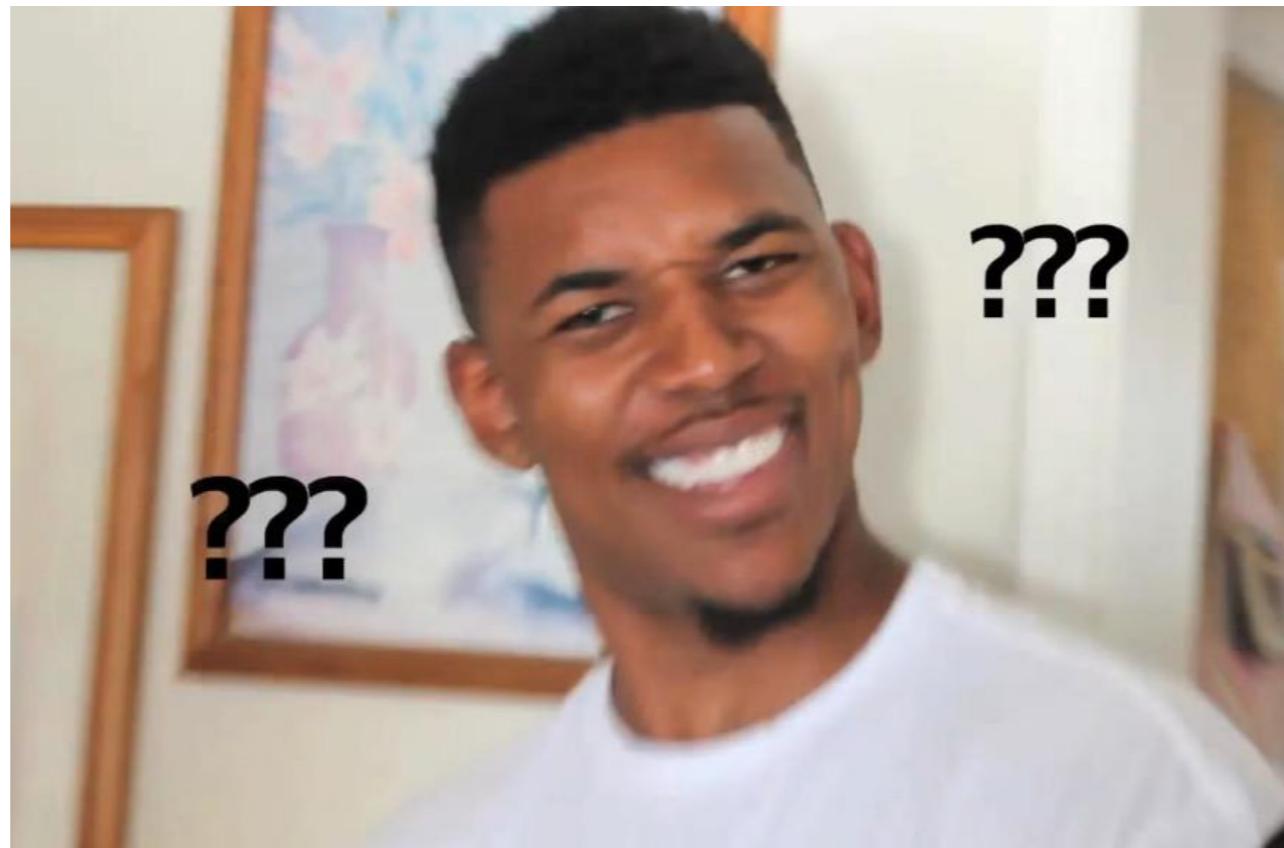
→ cosine similarity:

$$\begin{array}{cc} \img[alt="User 1 icon"]{} & \img[alt="User 2 icon"]{} \\ : & \end{array} \quad \frac{9 * 1 + 8.5 * 2}{23.96 * 17.46} = \textcolor{red}{0.062}$$

$$\begin{array}{cc} \img[alt="User 1 icon"]{} & \img[alt="User 3 icon"]{} \\ : & \end{array} \quad \frac{8.7 * 8 + 10 * 9 + 10 * 9 + 9 * 9 + 9 * 10 + 8 * 10 + 8.5 * 9.5}{23.96 * 24.94} = \textcolor{green}{0.973}$$

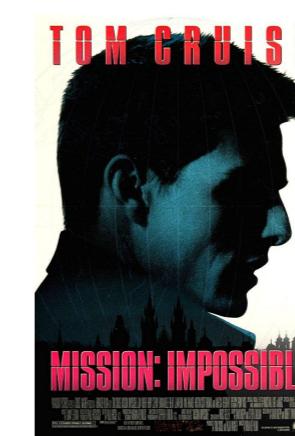
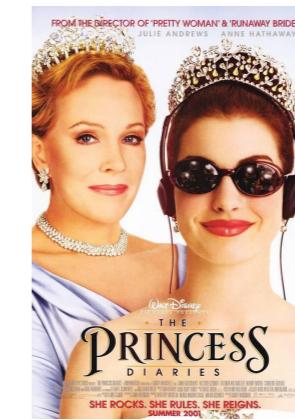
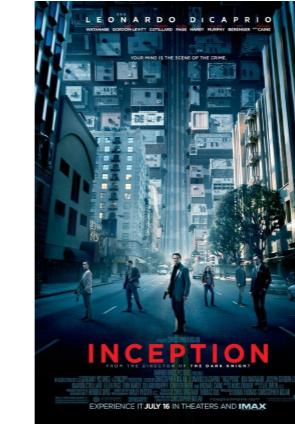
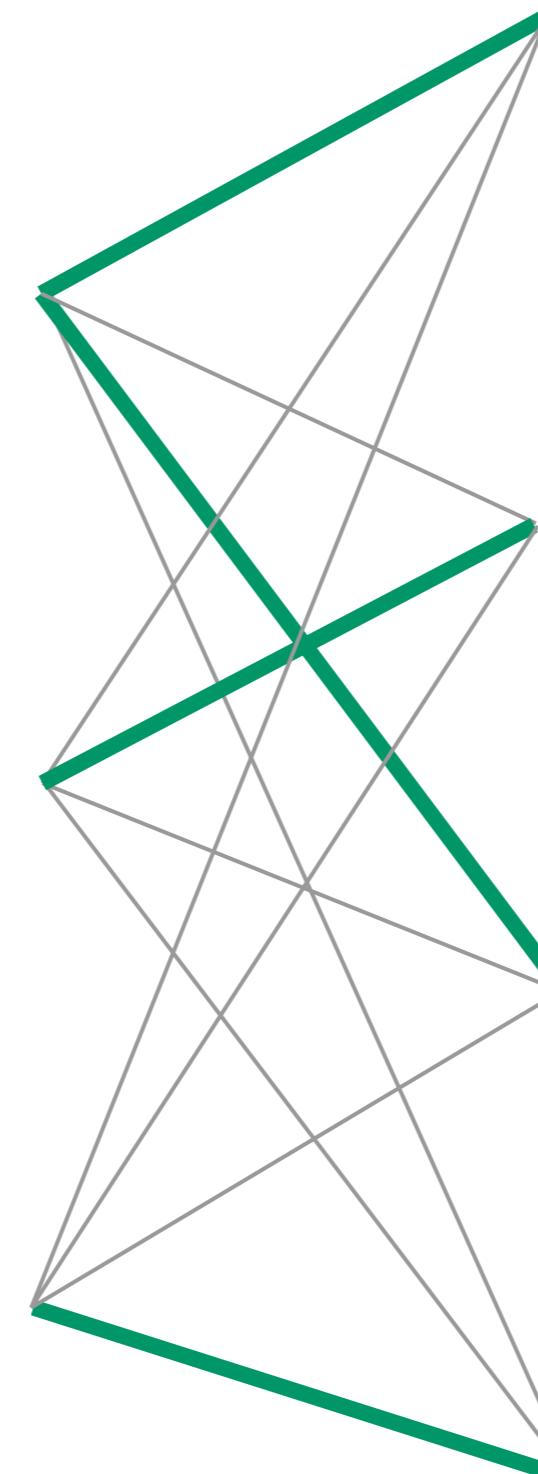
Latent Factor Models

- Explain the rating by characterizing both users and items on a number of *latent factors* inferred from rating patterns.



The factors are latent.

Not directly
observable ...



Not directly
observable ...

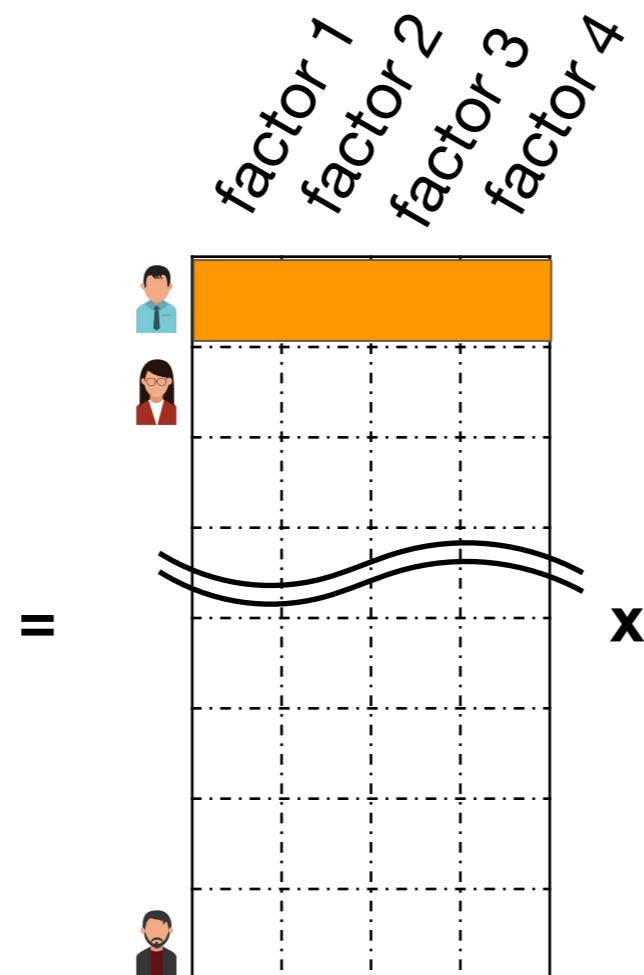


3	5	5	4	?	?	?	4	3	3
?	?	?	1	5	5	5	?	?	2
●									
●									
●									
●									
3	4	4	4	?	?	1	5	5	4



What is latent factor anyway?

	3	5	5	4	?	?	?	4	3	3
	?	?	?	1	5	5	5	?	?	2
	3	4	4	4	?	?	1	5	5	4



factor 1					
factor 2					
factor 3					
factor 4					

Singular Value Decomposition (SVD)

$$A = U \Sigma V^T$$

The diagram illustrates the Singular Value Decomposition (SVD) of a matrix A . It shows a 4x4 matrix A on the left, followed by an equals sign. To the right of the equals sign is a 4x4 matrix U , which is orthogonal. Next is a 4x4 diagonal matrix Σ , representing the singular value matrix. Following Σ is a 4x3 matrix V , which is also orthogonal. The matrices U , Σ , and V are labeled below their respective grid diagrams.

That is (kind of) what SVD is

3	5	5	4	?	?	?	4	3	3
?	?	?	1	5	5	5	?	?	2
3	4	4	4	?	?	1	5	5	4

=

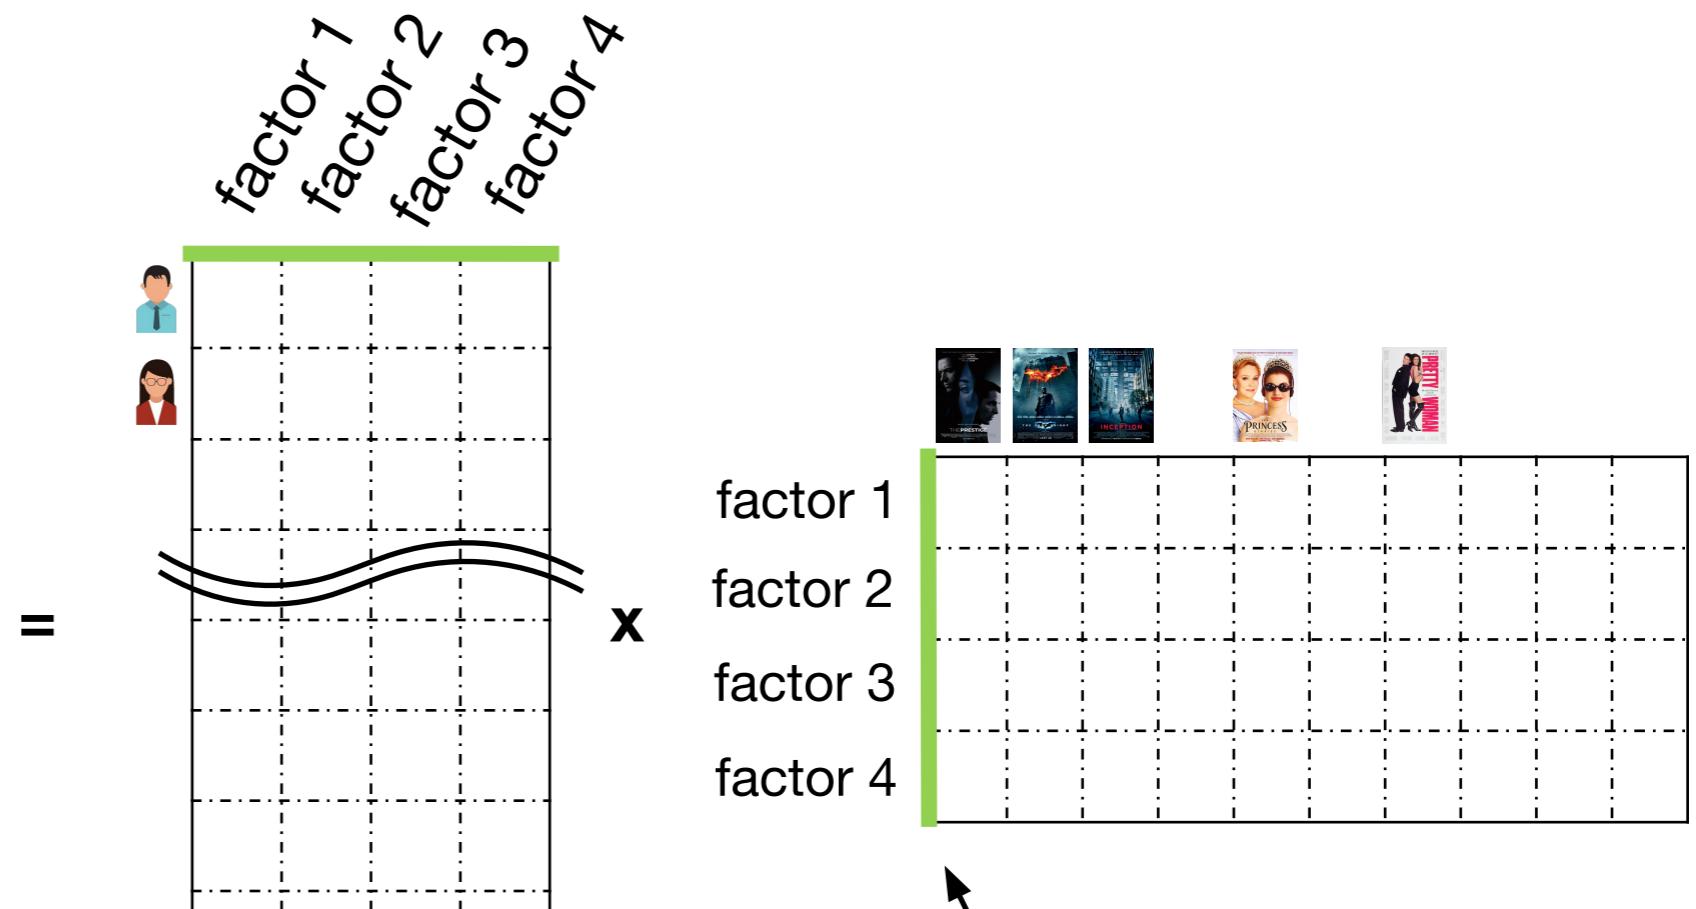
x

1	1	1	0	0	0	0	0	0	0
0	0	0	0	1	1	1	1	0	0
1	1	1	1	0	0	0	0	0	0
0	1	0	0	0	0	0	1	1	1

But in practice...

3	5	5	4	?	?	?	4	3	3
?	?	?	1	5	5	5	?	?	2
3	4	4	4	?	?	1	5	5	4

“Full of holes”

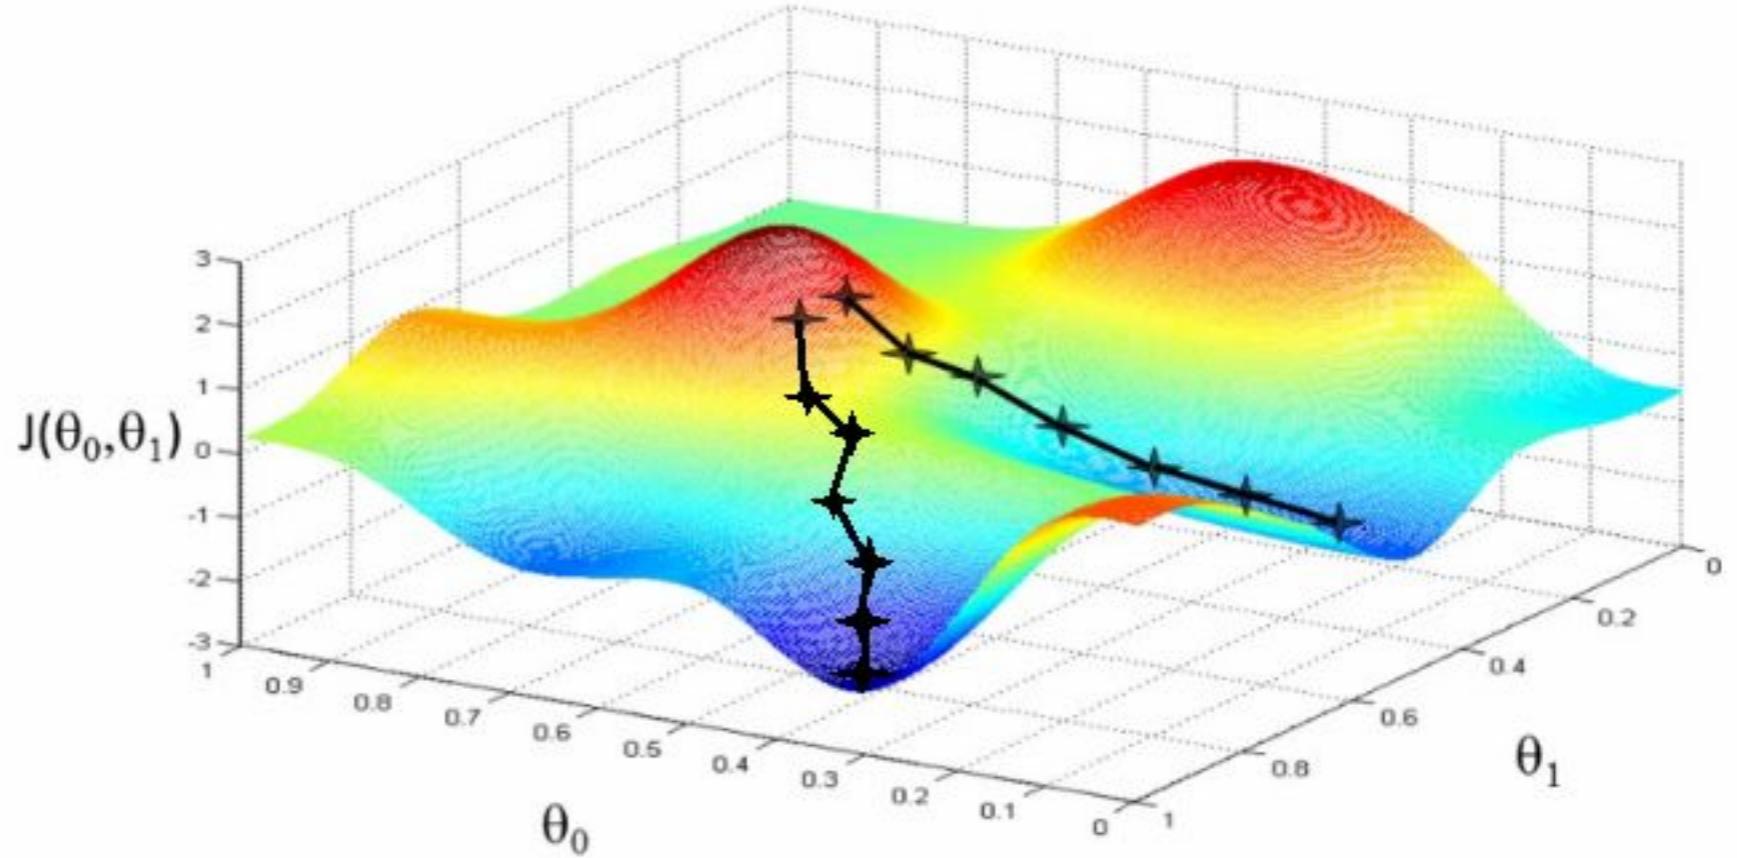
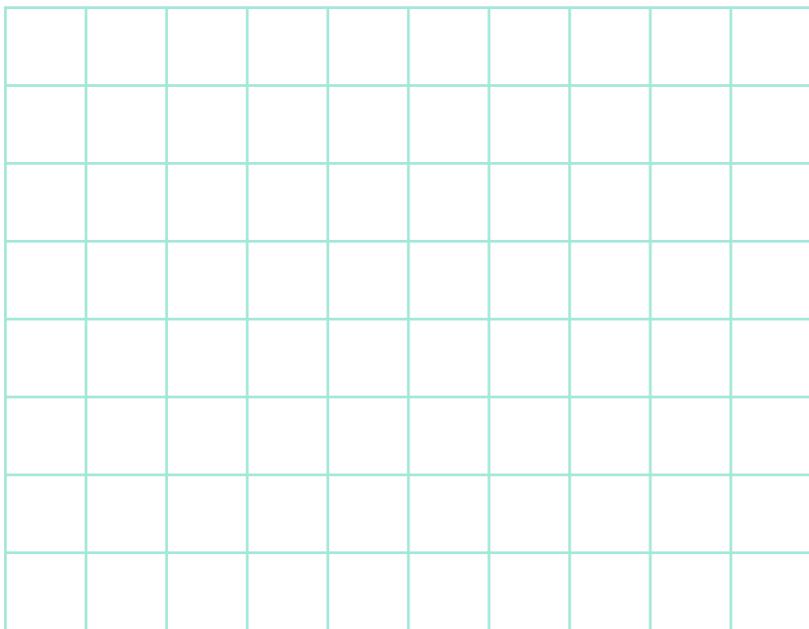


Can't change the length of “contact”

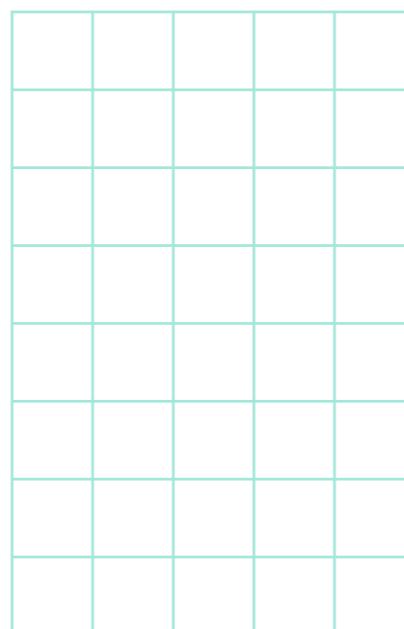
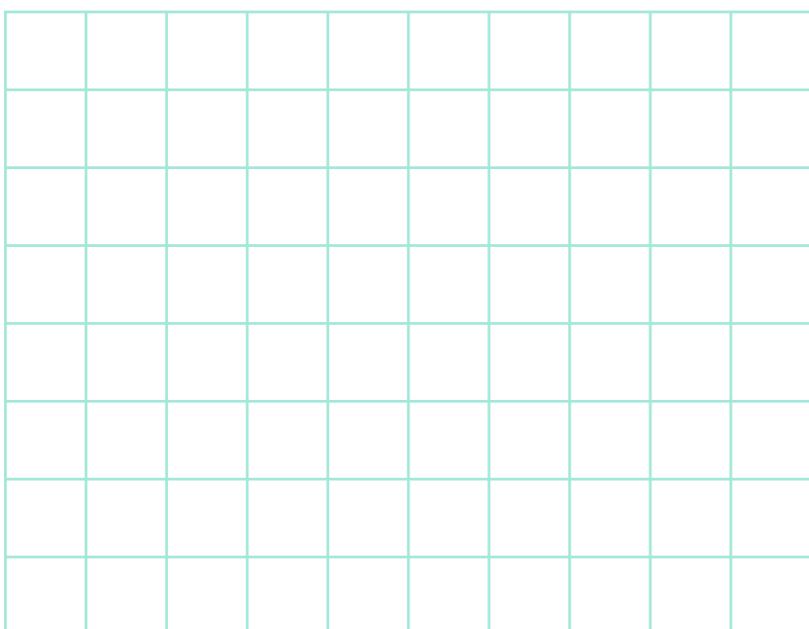
How to factor the matrix

- To perform matrix factorization for large matrices, we learn the entries through optimization methods such as *stochastic gradient descent (SGD)*.
- Methods like *alternating least square* (ALS) are also used when computation can be parallelized.
- We are going to briefly introduce SGD for its popularity.

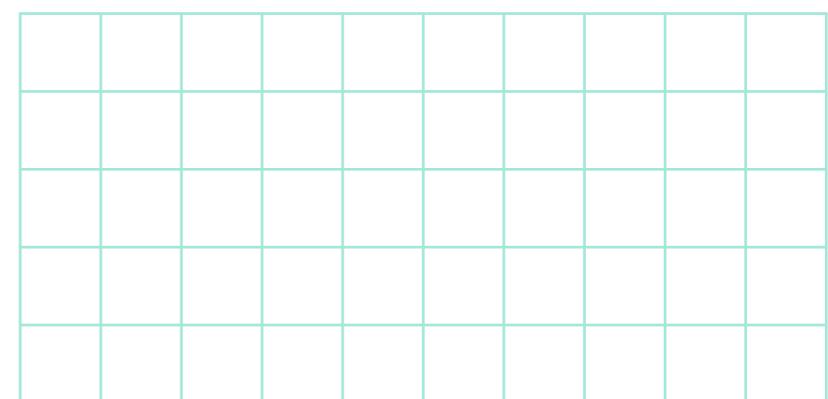
Goal:



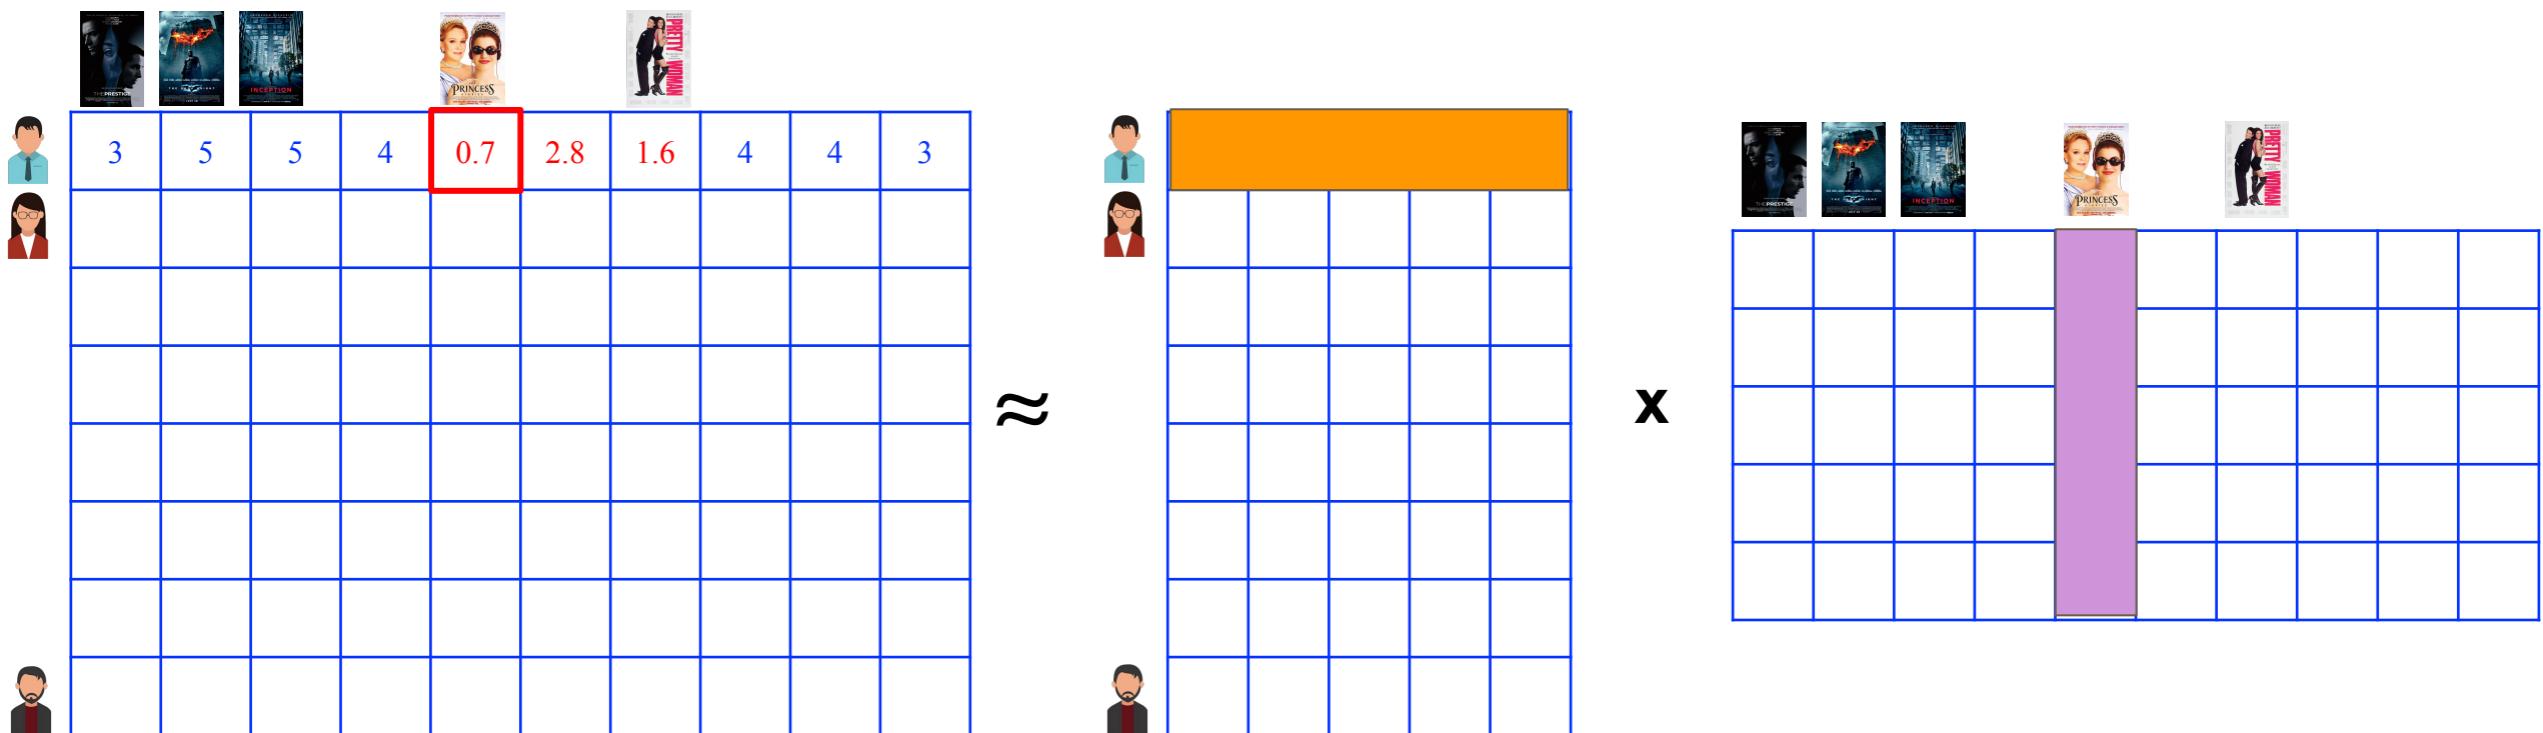
Initialize these and tune the numbers!



x



Ok, I've got the matrices. Then what?



→ Key points:

1. The approximate matrix will **not** be identical to the original.
2. The factor matrices will keep changing as long as there are users changing the rating (even if you stay inactive for a while).

Evaluating a model

- To evaluate how a machine learning model did, we use metrics such as *precision* and *recall*.

		Actual value	
		True	False
		True	False positive (FP)
Predicted value	True	True positive (TP)	
	False	False negative (FN)	True negative (TN)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Scenario 1:

		Actual value	
		Have cancer	Safe
Predicted value	Have cancer	45	5
	Safe	190	1000

$$\text{Precision} = \frac{45}{45 + 5} = 90\%$$

$$\text{Recall} = \frac{45}{45 + 190} = 19.15\%$$

Scenario 2:

		Actual value	
		Have cancer	Safe
Predicted value	Have cancer	200	800
	Safe	35	205

$$\text{Precision} = \frac{200}{200 + 800} = 20\%$$

$$\text{Recall} = \frac{200}{200 + 35} = 85.11\%$$

That's why we have F1-score

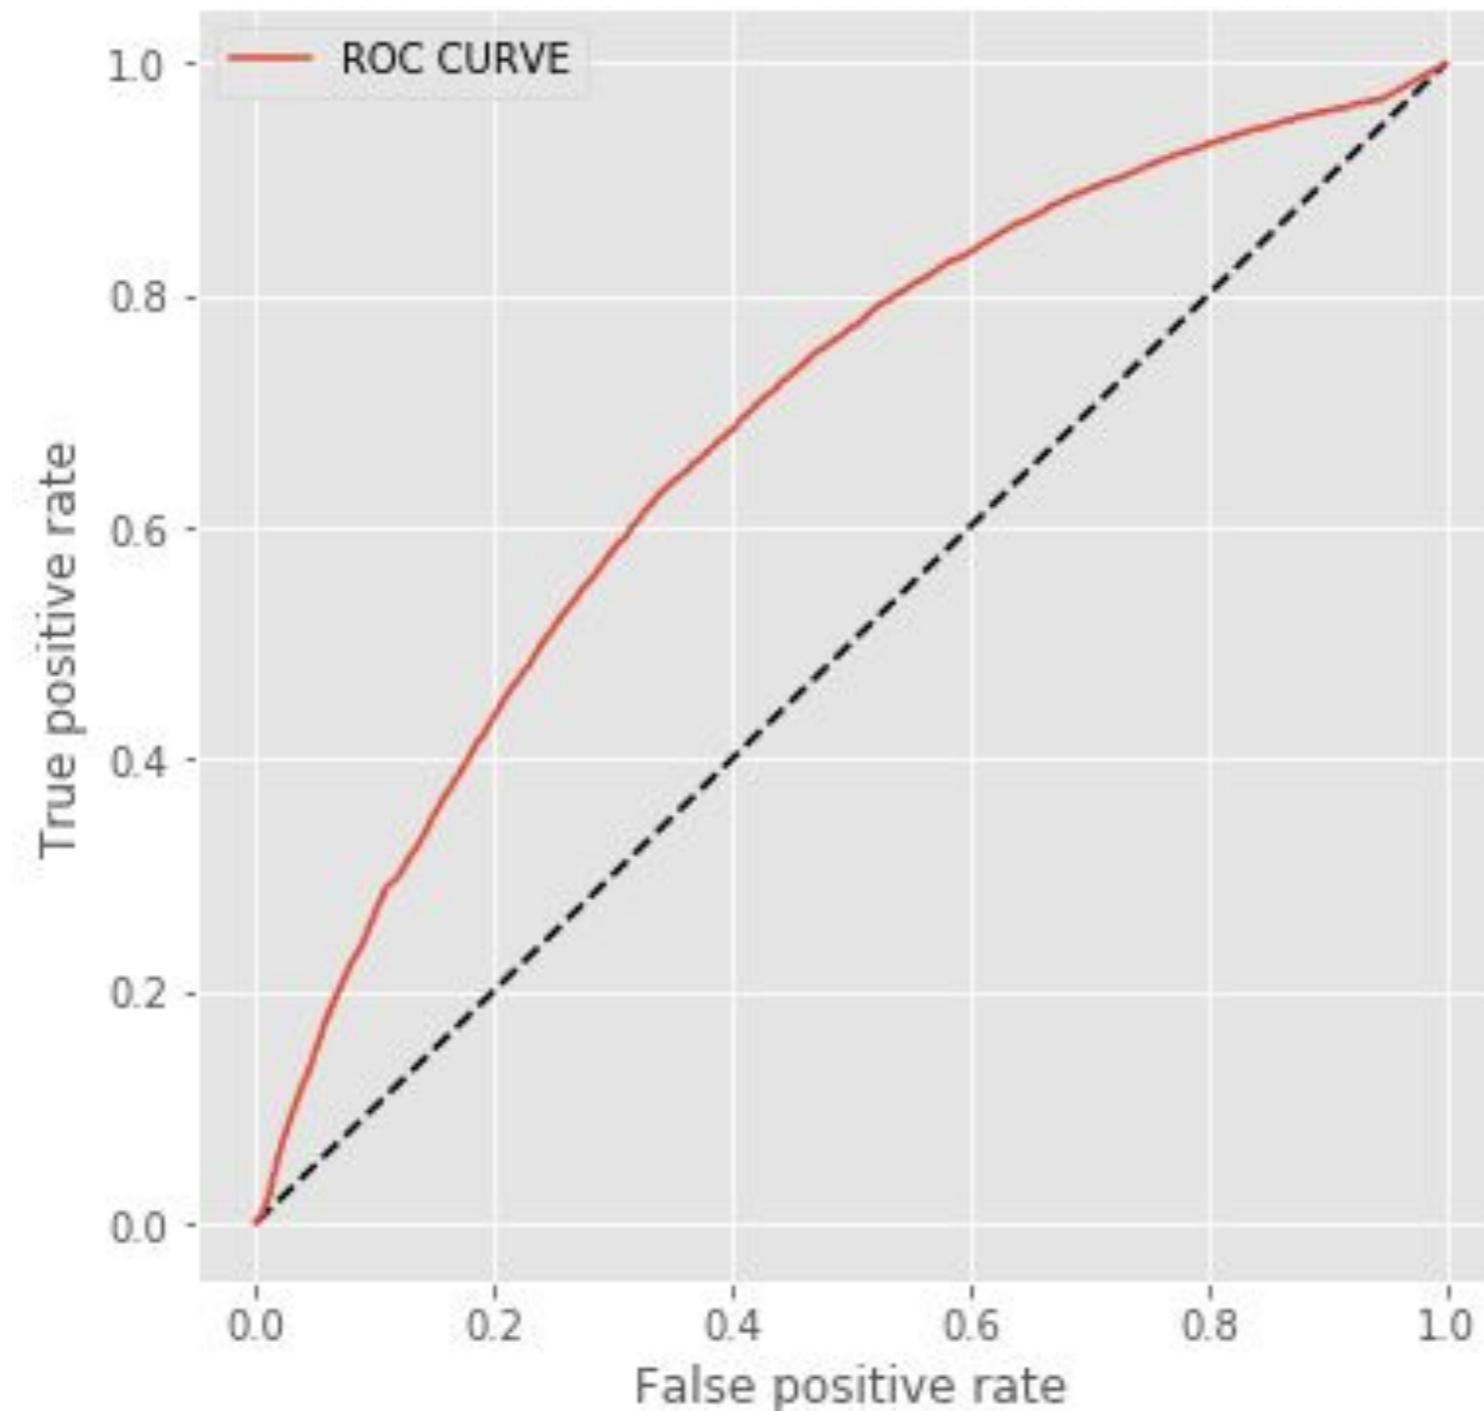
- F1-score (or F-score) is the harmonic mean between precision and recall:

$$F1 = 2 * \frac{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}}{2} = 2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

Also, the ROC curve



Receiver Operating Characteristic (ROC) curve



Let's evaluate the ranking as well

- A RE often delivers numerous outputs, while only a portion of them is most relevant to what the user really wants.
- We **rank our results** for the users, so the entries that the user would most likely selected would be near the top.
- How do we evaluate the *ranking* of results? 

鍾欣怡

Search



(鍾欣桐)



1



2



(鍾欣凌)

3

4

Relevance (0-3 scale):



1

Cumulative Gain: $1 + 3 + 3 + 1 = 10$
(@ rank 4)



3

Discounted Cumulative Gain (DCG): (@ rank p)

$$\sum_{i=1}^p \frac{rel_i}{\log_2(i+1)}$$

$$\frac{1}{\log_2(1+1)} + \frac{3}{\log_2(2+1)} + \frac{3}{\log_2(3+1)} + \frac{1}{\log_2(4+1)} = 4.824$$



3

$$\frac{3}{\log_2(1+1)} + \frac{3}{\log_2(2+1)} + \frac{1}{\log_2(3+1)} + \frac{1}{\log_2(4+1)} = 5.824$$



1

Normalized DCG:
(@ rank p)

$$\frac{DCG_p}{IDCG_p}, \quad IDCG_p = \sum_{i=1}^{|REL|} \frac{2^{rel_i} - 1}{\log_2(i+1)}$$